



Adaptive Computation Training Summary

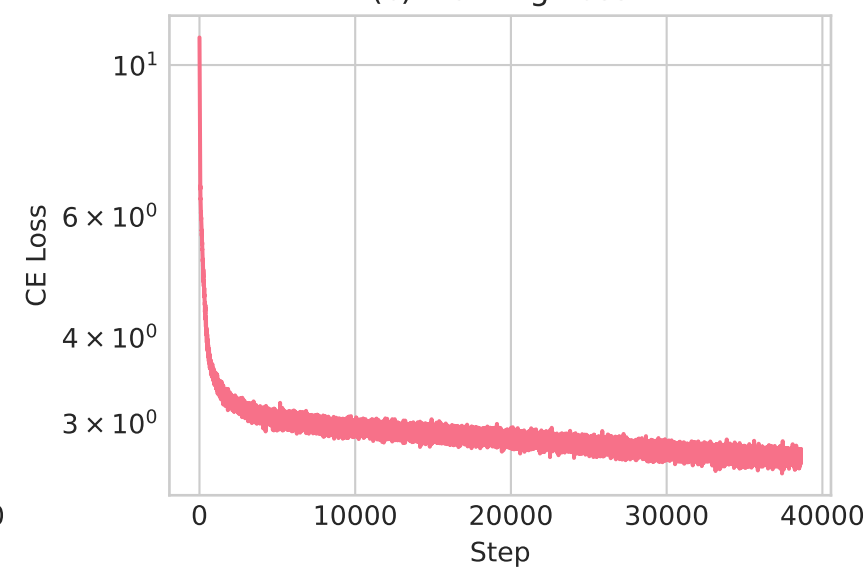
(a) Layer Specialization



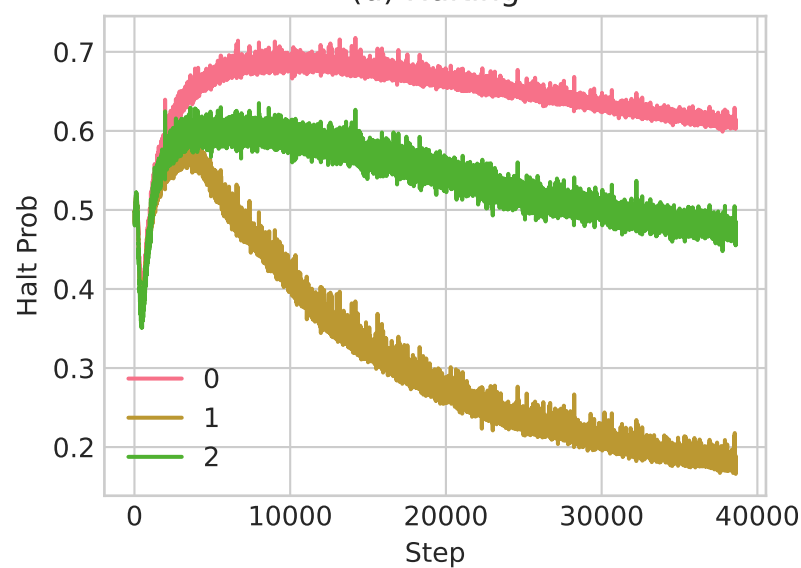
(b) Memory Gates



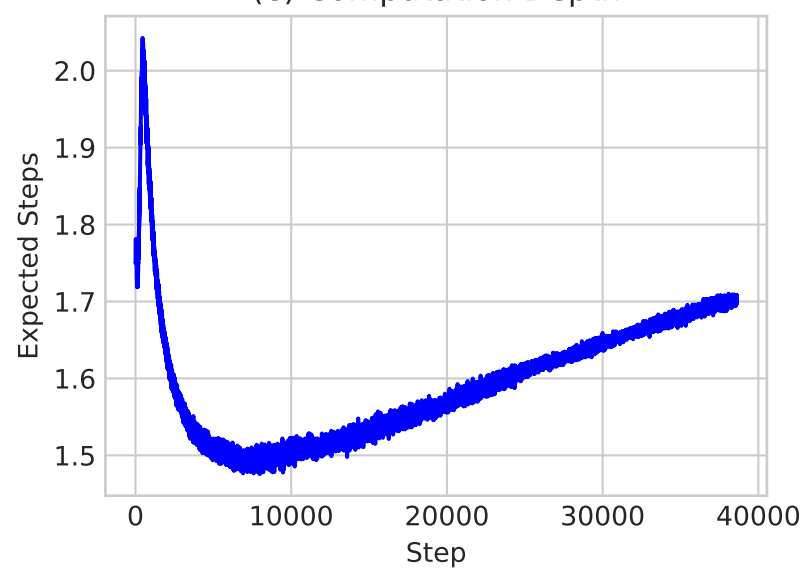
(c) Training Loss



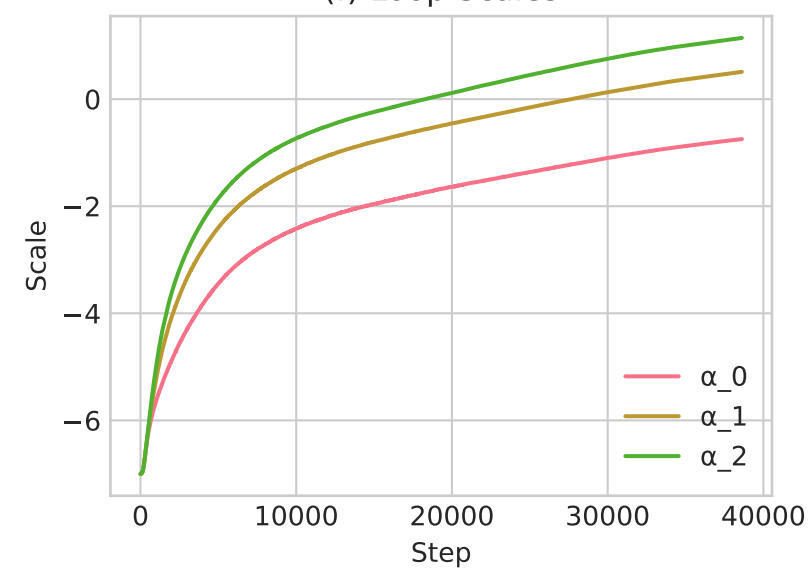
(d) Halting



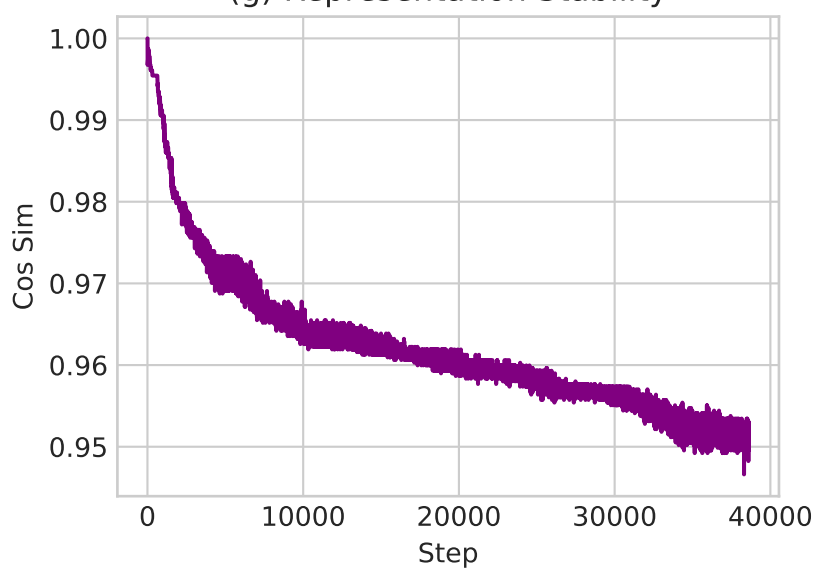
(e) Computation Depth



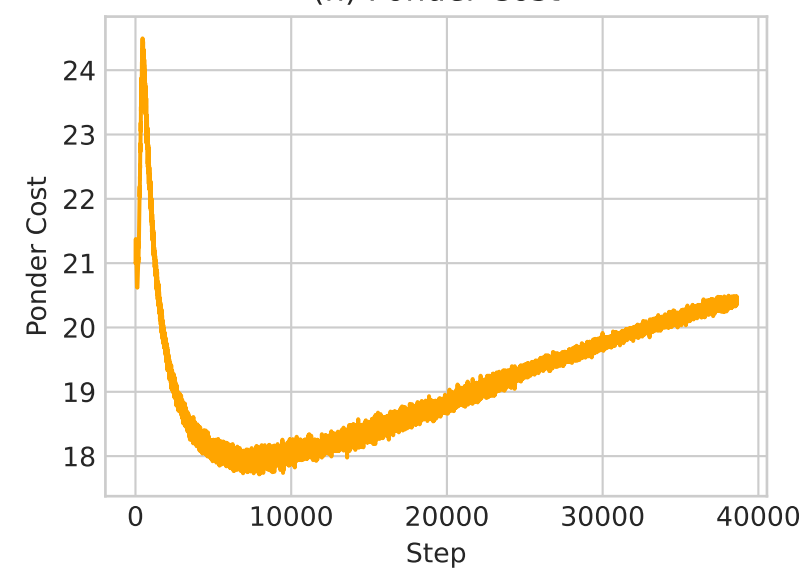
(f) Loop Scales



(g) Representation Stability



(h) Ponder Cost



(i) Gradients

